



CS 7140: ADVANCED MACHINE LEARNING

Maximum Likelihood Estimation

- Finding a hypothesis that fits the data well

$$\theta^{\star} = \operatorname{argmax}_{\theta} \log P(D | \theta, H)$$

- Work with the *logarithm* of the likelihood

- products of probabilities tends too be small

- likelihood multiples, log likelihood adds

- MLE is equivalent to minimize the relative entropy

$$KL(P(x | \theta^{\star}) || P(x | \theta)) = \mathbb{E}[\log P(x | \theta^{\star})] - H[P(x | \theta^{\star})]$$

Maximum a Posterior (MAP)

- $P(\theta | D) \propto P(D | \theta)P(\theta)$

posterior likelihood prior

- Conjugate distributions

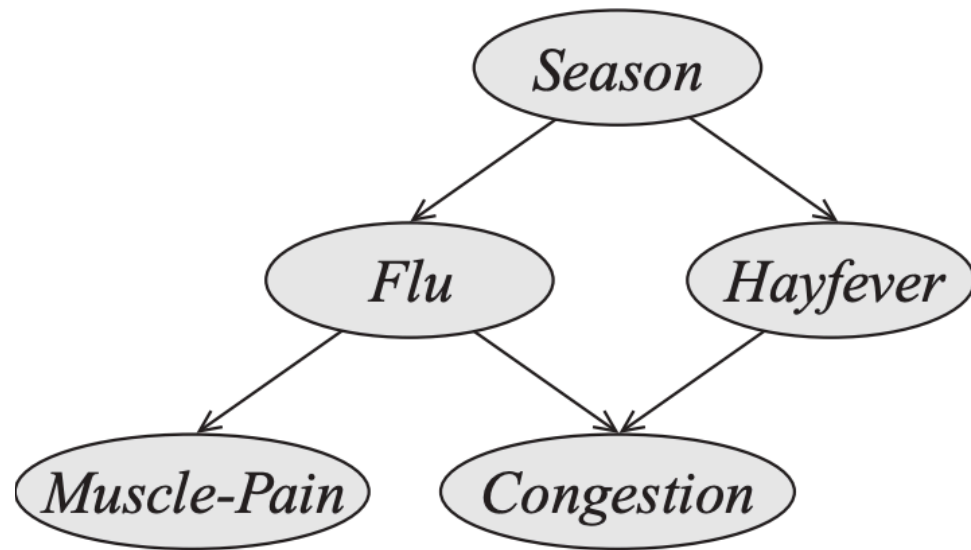
$$P(\theta | D) \propto P(D | \theta)P(\theta)$$


same distribution family

- Example: Dirichlet prior is conjugate to multinomial
if $P(\theta) = \text{Dir}(\alpha_1, \dots, \alpha_K)$ then
 $P(\theta | D) = \text{Dir}(M_1 + \alpha_1, \dots, M_K + \alpha_K)$

STRUCTURE LEARNING

Structure Learning



- Reconstruct the structure of Bayesian networks G^\star
- Find independencies in variables
- G^\star is **not** identifiable: the class of all I-equivalent networks

I-Equivalence


- Two Bayesian networks are I-equivalent if they encode precisely the same conditional independence assertions.
- Are these I-equivalent?




(a)





(b)



(c)



(d)



Two Bayesian networks are I-equivalent if they have the same **skeletons** and same **V-structures**.

Methods Overview


- **Constraint-based** : reconstruct a network structure that best captures the independencies
- **Score-based**: find a network structure that best fits the observed data
- **Bayesian model averaging**: average the prediction of all possible structures.

Constraint-Based Structure Learning

Learn an I-Equivalent class and find minimal I-Map


Algorithm input X_1, \dots, X_n , output G^\star

1. Build a complete graph for X_1, \dots, X_n
2. For every pair X_i, X_j , find the witness set U such that $X_i \perp X_j \mid U$
3. Remove edge $X_i - X_j$ if U is not empty
4. Identify triplets that are *immoralities*
5. Orient edges by propagating a set of *constraints*



Identify Immoralities (V-Structures)


- for every triplets that satisfies $X_i - X_j - X_k$ but $X_i \neg X_k$
- Find the witness set of X_i, X_k such that $X_i \perp X_k \mid U_{X_i, X_k}$
- If $X_j \notin U_{X_i, X_k}$, add $X_i \rightarrow X_j$ and $X_j \leftarrow X_k$



Independence Test hypothesis testing $P(X_1, X_2) = P(X_1)P(X_2)$

Rules for Orienting Edges

- Preserve the V-structures
- Acyclic graph
- Allow undirected edges to have both directions



Guaranteed Recovery

The algorithm reconstructs the network with *polynomial* number of *independence tests*

Assumptions


- The network G^\star has bounded indegree d
- The independence procedure can perfectly answer any independence query that involves up to $2d + 2$ variables.
- The underlying distribution P^\star is faithful to G^\star

Score-Based Structure Learning

Optimize the network structure score that best fits the data

Algorithm input X_1, \dots, X_n , output G^\star , a scoring function

1. Generate a set of possible network structures with bounded indegree
2. Search the space of graphs
3. Return the high-scoring one

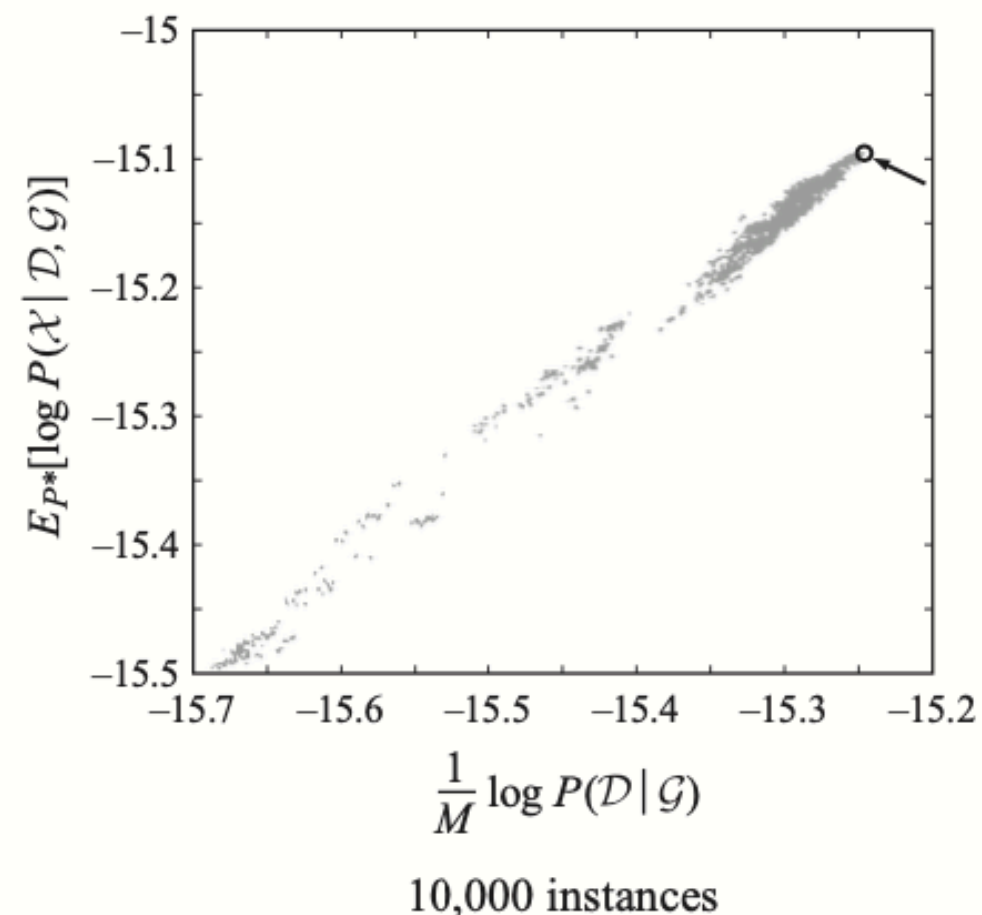



Maximum Likelihood Score

Find the model that has the highest likelihood

$$\max_G \max_{\theta_G} L(D | G, \theta_G)$$

- Never prefers simpler network
- Add an edge will never decrease the score, overfit the training data



Bayesian Score

Place a prior whenever we have uncertainty

$$P(G|D)$$


- By Bayesian rule: $P(G|D) = \frac{P(D|G)P(G)}{P(D)}$

- Score: $\log P(D|G) + \log P(G)$

marginal likelihood


- Intuition:

$$\frac{1}{M} \log P(D|G) \approx \mathbb{E}_P[\log P(X|G, D)]$$



Bayesian Information Criteria

- Biased toward simpler structures, willing to recognize more complex structures with more data.
- Trade off the likelihood — fit to data — and some notion of model complexity , reducing the extent of *overfitting*.
- $\text{score}_{BIC}(G : D) = L(D | \theta) - \frac{\log M}{2} \text{Dim}(G)$










Score Decomposition

- Specify prior over structure $P(G) \propto \prod_i P(\text{Pa}_{X_i} = \text{Pa}_{X_i}^G)$
- Decomposable score: $\text{score}(G : D) = \sum_i \text{score}(X_i | \text{Pa}_{X_i}^G : D)$
- A local change in the structure (such as adding an edge) does not change the score of other parts of the structure that remained the same.
- Reduce dramatically the computational overhead of evaluating different structures during search.

Structure Search

- **State:** a network structure of variables
- **Action:** edge addition, edge deletion, edge reversal



Bayesian Model Averaging

Compute the average prediction over candidate network structures

$$\mathbb{E}_{P(G|D)}[f(G)] = \sum_G f(G)P(G|D)$$


Algorithm input X_1, \dots, X_n , output G^\star , a prediction function

1. Find a set \mathcal{G}' of high scoring structures

2. Estimate the probability of the structure in \mathcal{G}'

3. Compute average prediction


$$P(f|D) \approx \frac{\sum_{G \in \mathcal{G}'} P(G|D)f(G)}{\sum_{G \in \mathcal{G}'} P(G|D)}$$




$$f(G_1) + f(G_2)$$

Find Candidate Structures


- Large data: a single high-scoring structure gives a good approximation of the prediction
- Small data: a large number of high-scoring models
- Alternative: MCMC over structures



(a)



(b)




(c)

MCMC over Structures


- Use Metropolis-Hasting as an example
- The current state is G , sample G' from the proposal distribution
- Accept the transition with probability

uniform distribution


$$\min \left[1, \frac{P(G', D) T^Q(G' \rightarrow G)}{P(G, D) T^Q(G \rightarrow G')} \right]$$



(a)



(b)



(c)